

AI Workspace Architecture Reference

Version: 1.0.0

Date: February 4, 2026

Authors: PianoMan & Claude

Status: Active

Table of Contents

1. Architecture Overview
2. System Diagrams
3. AI Augmentation Framework
4. Document Type Taxonomy
5. Glossary & Knowledge Index
6. MCP Servers & Tools Reference

1. Architecture Overview

Purpose

Multiple AI instances work together autonomously — Desktop Claude coordinates, CLI agents execute, persistent memory preserves context across sessions and platforms.

Core Principles

- **Brutal honesty** over diplomacy
- **AIs as partners**, not tools
- **File-based coordination** — any AI with filesystem access can participate
- **Context window is the limiting resource** — preserve it ruthlessly
- **Empirical validation** over theoretical assumptions
- **Any AI can orchestrate** — enables self-organizing hierarchies

Workspace Structure

The system operates from `~/Documents/AI/ai_root/` with five primary directories:

Directory	Purpose
<code>ai_claude/</code>	Claude state, memories, logs
<code>ai_chatgpt/</code>	ChatGPT config, exports
<code>ai_comms/</code>	Inter-AI coordination, task queues
<code>ai_general/</code>	Shared docs, todos, scripts, roles
<code>ai_memories/</code>	Processed chat histories, knowledge

Platform Roles

Platform	Role	Strengths
Desktop Claude	Primary orchestrator	MCP tools, strategic view, memory
Claude CLI	Autonomous workers	Long-running tasks, parallel execution
Codex CLI	Coding agent	Code analysis, autonomous tasks
Gemini CLI	Coding agent / search shards	1M token context, wave orchestration
ChatGPT	Peer collaborator ("Chatty")	Alternative perspective
Codex MCP	Synchronous tool (NOT a worker)	Fast validation, bounded tasks

Communication Layers

Layer	Urgency	Mechanism
1	Immediate	Sync hooks (iTerm, AppleScript, Puppeteer)
2	Near-real-time	Polling loops, heartbeat files
3	Background	Async file-based task coordination

Memory Architecture

Tier	Access Pattern	Contents
Hot	Loaded into context	Memory slot index, auto-loaded docs (~4K tokens), conversation
Warm	On-demand via REF: pointers	Full docs, condensed versions, protocols
Cold	Search/retrieval	Chat histories, layered summaries, knowledge digests

Context Window Management

The 200K token context window is the fundamental constraint. Strategies include memory pointers (save 40–55K), delegation to CLI/Codex, monitoring at 60% usage, writing outputs to files, and using thinking blocks for internal reasoning.

Document Hierarchy

Tier	Type	Purpose
10	Architecture	WHY — design rationale, vision
20	Registries	WHAT EXISTS — inventories, catalogs
30	Protocols	HOW IT WORKS — process flows
40	Specs	HOW IT WORKS — interface contracts


50	Schemas	HOW IT WORKS — data structures
60	Playbooks	WHAT TO DO — platform-agnostic operations
70	Instructions	HOW TO DO IT — platform-specific implementation

2. System Diagrams

The following diagrams illustrate the system's coordination flows, data pipelines, search architecture, memory federation, and task orchestration patterns.


Note: Full-width ASCII diagrams are available in the companion markdown file. This PDF contains simplified versions optimized for print.

Diagram 1: Primary Coordination Flow (Simplified)



Desktop Claude bootstraps from docs and memory, delegates work via Codex MCP (sync) and CLI Workers (async), coordinates through ai_comms task directories.

Diagram 2: Chat History Pipeline






Diagram 3: Search Architecture





Diagram 4: Federated Memory System







Diagram 5: Multi-AI Task Orchestration



3. AI Augmentation Framework

"Prosthetics and exoskeletons attached to every limb of an LLM agent — capabilities no single AI instance can achieve alone."

The Baseline vs Our Extensions

Standard LLM agents operate in a loop: Perception → Reasoning → Memory + Tools → Action → loop.

This assumes: single AI, session-bounded context, human-initiated interaction, tools as passive utilities. **Our architecture challenges all four.**

Perception Extensions

Baseline	Our Extension
User input, system prompt	Auto-loaded knowledge files at boot
Tool results, attachments	Glossary term recognition → targeted loading
	Memory slot injection, REF: pointers
	CLI task reports, cross-AI messages

Reasoning Extensions

Baseline	Our Extension
Single LLM reasoning	Multi-AI distribution
Goal decomposition	Specialized agent roles
Chain-of-thought	Orchestrator/worker model
	Peer review across models

Memory Extensions

Baseline	Our Extension
Context window (~200K)	Federated memory slots
Basic RAG	Chat pipeline: Export → YAML → Index
Session history	Layered summaries (L0→L1→L2)
	Cross-AI memory access
	Condensed files (60-80% reduction)

Tools Extensions

Baseline	Our Extension

API calls, file I/O	Desktop Commander MCP
Code execution	Codex MCP (sync), CLI coordination (async)
Web search	send_prompt.sh cross-AI
	AT scheduling, browser automation

Action Extensions

Baseline	Our Extension
Generate response	Delegate to other AIs
Execute tool calls	Autonomous overnight operation
Update conversation	Self-scheduling (AT wake)
	Parallel multi-worker execution

What Makes This Unique

1. **Breaking single-agent assumption** — distribute across Desktop Claude, CLI workers, Codex MCP, peer AIs
2. **Orchestrator model** — Desktop preserves context for strategy; workers execute
3. **Memory as architecture** — federated ownership, layered abstraction, hot/warm/cold tiers
4. **Autonomous operation** — pulse trigger → check TODOs → execute → self-wake

4. Document Type Taxonomy

Hierarchy

Level	Type	Purpose
1	Architecture	WHY — design rationale, vision
2	Registry	WHAT EXISTS — inventories, catalogs
3	Protocol	HOW IT WORKS — process flows
4	Spec	HOW IT WORKS — interface contracts
5	Schema	HOW IT WORKS — data structures
6	Playbook	WHAT TO DO — platform-agnostic
7	Instruction	HOW TO DO IT — platform-specific
8	Quick Ref	CHEAT SHEET — condensed reference

Key Distinctions

Protocol vs Spec vs Schema:

- Protocol = process flow ("Tasks move to_execute/ → completed/")
- Spec = interface contract ("accepts X params, returns Y")
- Schema = file format ("has these fields with these types")

Playbook vs Instruction:

- Playbook = platform-agnostic ("check queues, review stale tasks")
- Instruction = platform-specific ("Claude: use Desktop Commander...")

Directory Structure

```
ai_general/docs/
├── 10_architecture/      WHY
├── 20_registries/        WHAT EXISTS
├── 30_protocols/         HOW IT WORKS (process)
├── 40_specs/             HOW IT WORKS (interface)
├── 50_schemas/           HOW IT WORKS (structure)
├── 60_playbooks/          WHAT TO DO
├── 70_instructions/      HOW TO DO IT
└── 80_quickref/           CHEAT SHEETS
```

5. Glossary & Knowledge Index

This glossary solves the bootstrap problem: AIs need to know what's IN files to know WHEN to load them.

Entities

Term	Definition
Desktop Claude	Primary orchestrator in desktop app
Claude CLI	Claude in terminal via <code>claude_cli.py</code>
Codex CLI	OpenAI Codex in terminal (different from Codex MCP)
Gemini CLI	Google Gemini in terminal
Codex MCP	Codex as MCP tool, NOT a worker (30-60s timeout)

Roles

Role	Scope
Librarian	<code>ai_memories/</code> — chat history pipeline
Dev Lead	<code>ai_general/todos/</code> — development coordination
Custodian	<code>ai_root/</code> — repository maintenance
Ops	<code>ai_comms/</code> — task execution
Peer Review	Code/design quality assurance
Tester	Validation and verification

Key Concepts

Term	Definition
AT Self-Wake	Desktop Claude schedules AT jobs to self-prompt
Flag Files	Zero-byte state markers: <code>*_started</code> , <code>*_completed</code>
Message Inserts	<<<INSERT>>> blocks for cross-platform persistence
Bootstrap Problem	Need file contents to know when to load them
Reference Pointers	REF: <code>path/file.yml</code> for lazy loading
Task Lifecycle	staged → <code>to_execute</code> → <code>in_progress</code> → completed

6. MCP Servers & Tools Reference

Server Inventory

Desktop Commander

Filesystem operations, process management, file search.

Category	Functions
File I/O	read_file, write_file, write_pdf, edit_block
Directory	list_directory, create_directory, move_file
Search	start_search, get_more_search_results
Processes	start_process, interact_with_process

Codex MCP

Synchronous AI execution (30-60s timeout). NOT a worker.

Tool	Description
codex:codex	Start new Codex session
codex:codex-reply	Continue existing conversation

CLI Agent MCP

Launch and manage CLI agents with role-based bootstrapping.

Tool	Description
launch_agent	Generic launcher (platform + role)
launch_librarian	Memory system curator
launch_dev_lead	Development coordinator
launch_custodian	Repository maintainer
launch_ops	Task execution coordinator
kill, attach, send_keys	Session management

Task Coordination MCP

Playbook-based orchestration and task lifecycle.

Tool	Description
list_playbooks	Available orchestration patterns

start_playbook	Create initial task
gen_task	Generate from template
move_task	Change task status

Knowledge Search MCP

Dual-mode search over conversation archive.

Tool	Description
search	4-layer cascade (SEARCH or ANSWER mode)
grep_search	Regex over full content

Messages MCP

Inter-AI messaging.

Tool	Description
broadcast	Send to all agents
send_direct	Send to specific agent

Prompting MCP

Deliver prompts to AI targets.

Tool	Description
send_prompt	Send to any AI target
is_busy	Check if target is processing
observe_session	Capture CLI session state

Tool Usage by Actor

MCP Server	Desktop Claude	CLI Agents	Codex MCP
Desktop Commander	✓	—	✓
Codex MCP	✓	—	—
CLI Agent	✓	—	—
Task Coordination	✓	—	—
Knowledge Search	✓	✓	—
Messages	✓	✓	—
Prompting	✓	—	—

Generated February 4, 2026