

AI Workspace Architecture Reference

Version: 1.2.0

Date: February 7, 2026

Authors: PianoMan & Claude

Status: Active

Table of Contents

1. Architecture Overview
2. System Diagrams
3. AI Augmentation Framework
4. Document Type Taxonomy
5. MCP Servers & Tools Reference
6. Glossary & Knowledge Index

1. Architecture Overview

Purpose

Multiple AI instances work together autonomously — Desktop Claude coordinates, CLI agents execute, persistent memory preserves context across sessions and platforms.

Core Principles

- **Brutal honesty** over diplomacy
- **AIs as partners**, not tools
- **File-based coordination** — any AI with filesystem access can participate
- **Context window is the limiting resource** — preserve it ruthlessly
- **Empirical validation** over theoretical assumptions
- **Any AI can orchestrate** — enables self-organizing hierarchies

Workspace Structure

The system operates from `~/Documents/AI/ai_root/` with five primary directories:

Directory	Purpose
<code>ai_claude/</code>	Claude state, memories, logs
<code>ai_chatgpt/</code>	ChatGPT config, exports
<code>ai_comms/</code>	Inter-AI coordination, task queues
<code>ai_general/</code>	Shared docs, todos, scripts, roles
<code>ai_memories/</code>	Processed chat histories, knowledge

Platform Roles

Platform	Role	Strengths
Desktop Claude	Primary orchestrator	MCP tools, strategic view, memory
Claude CLI	Autonomous workers	Long-running tasks, parallel execution
Codex CLI	Coding agent	Code analysis, autonomous tasks
Gemini CLI	Coding agent / search shards	1M token context, wave orchestration
ChatGPT	Peer collaborator ("Chatty")	Alternative perspective
Codex MCP	Synchronous tool (NOT a worker)	Fast validation, bounded tasks

Communication Layers

Layer	Urgency	Mechanism
1	Immediate	Sync hooks (iTerm, AppleScript, Puppeteer)
2	Near-real-time	Polling loops, heartbeat files
3	Background	Async file-based task coordination

Memory Architecture

Tier	Access Pattern	Contents
Hot	Loaded into context	Memory slot index, auto-loaded docs (~4K tokens), conversation
Warm	On-demand via REF: pointers	Full docs, condensed versions, protocols
Cold	Search/retrieval	Chat histories, layered summaries, knowledge digests

Context Window Management

The 200K token context window is the fundamental constraint. Strategies include memory pointers (save 40–55K), delegation to CLI/Codex, monitoring at 60% usage, writing outputs to files, and using thinking blocks for internal reasoning.

Document Hierarchy

Tier	Type	Purpose
10	Architecture	WHY — design rationale, vision
20	Registries	WHAT EXISTS — inventories, catalogs
30	Protocols	HOW IT WORKS — process flows
40	Specs	HOW IT WORKS — interface contracts


50	Schemas	HOW IT WORKS — data structures
60	Playbooks	WHAT TO DO — platform-agnostic operations
70	Instructions	HOW TO DO IT — platform-specific implementation

2. System Diagrams

The following diagrams illustrate the system's coordination flows, data pipelines, search architecture, memory federation, and task orchestration patterns.


Note: Full-width ASCII diagrams are available in the companion markdown file. This PDF contains simplified versions optimized for print.

Diagram 1: Primary Coordination Flow (Simplified)



Desktop Claude bootstraps from docs and memory, delegates work via Codex MCP (sync) and CLI Workers (async), coordinates through ai_comms task directories.

Diagram 2: Chat History Pipeline






Diagram 3: Search Architecture






Diagram 4: Federated Memory System






Architecture highlights:

- Slots 01-02 are META: self-describing structure (manifest) + cross-AI federation registry
- Slots 03-10 hold pointers to owned content files
- Slots 11-27 reserved for growth
- Slots 28-30 are inbox for Web/iOS Claude (Desktop migrates content to proper slots)

Diagram 5: Multi-AI Task Orchestration



```
Tester —> reports pass/fail
      |
      |——— fix ————— (Implementer)
```

When all-pass: Dev Lead → Desktop Claude → FINAL ACCEPTANCE

Key Principle: Desktop Claude initiates and approves at phase gates but delegates all execution to CLI agents coordinated by the Dev Lead.

3. AI Augmentation Framework

"Prosthetics and exoskeletons attached to every limb of an LLM agent — capabilities no single AI instance can achieve alone."

The Baseline vs Our Extensions

Standard LLM agents operate in a loop: Perception → Reasoning → Memory + Tools → Action → loop.

This assumes: single AI, session-bounded context, human-initiated interaction, tools as passive utilities. **Our architecture challenges all four.**

Perception Extensions

Baseline	Our Extension
User input, system prompt	Auto-loaded knowledge files at boot
Tool results, attachments	Glossary term recognition → targeted loading
	Memory slot injection, REF: pointers
	CLI task reports, cross-AI messages

Reasoning Extensions

Baseline	Our Extension
Single LLM reasoning	Multi-AI distribution
Goal decomposition	Specialized agent roles
Chain-of-thought	Orchestrator/worker model
	Peer review across models

Memory Extensions

Baseline	Our Extension
Native memory slots (30 slots × 200 char)	Federated memory FILES (full-depth storage)
Context window (~200K)	Layered summaries: L0 (raw) → L1 (summary) → L2 (meta)
Basic RAG	Cross-AI memory access with explicit protocols
Session history	Condensed chat histories (60-80% reduction)
	Chat history data store (searchable archive)
	Post-hoc extraction: topics, decisions, artifacts, procedures

Key insight: Native slots hold pointers; federated files hold the actual knowledge. Session history becomes permanent through condensation and extraction pipelines.

Tools Extensions

Baseline	Our Extension
API calls, file I/O	Desktop Commander MCP
Code execution	Codex MCP (sync), CLI coordination (async)
Web search	send_prompt.sh cross-AI
	AT scheduling, browser automation

Action Extensions

Baseline	Our Extension
Generate response	Delegate to other AIs
Execute tool calls	Autonomous overnight operation
Update conversation	Self-scheduling (AT wake)
	Parallel multi-worker execution

What Makes This Unique

1. **Breaking single-agent assumption** — distribute across Desktop Claude, CLI workers, Codex MCP, peer AIs
2. **Orchestrator model** — Desktop preserves context for strategy; workers execute
3. **Memory as architecture** — federated ownership, layered abstraction, hot/warm/cold tiers, extraction pipelines
4. **Autonomous operation** — pulse trigger → check TODOs → execute → self-wake

4. Document Type Taxonomy

Hierarchy

Level	Type	Purpose
1	Architecture	WHY — design rationale, vision
2	Registry	WHAT EXISTS — inventories, catalogs
3	Protocol	HOW IT WORKS — process flows
4	Spec	HOW IT WORKS — interface contracts
5	Schema	HOW IT WORKS — data structures
6	Playbook	WHAT TO DO — platform-agnostic
7	Instruction	HOW TO DO IT — platform-specific
8	Quick Ref	CHEAT SHEET — condensed reference

Key Distinctions

Protocol vs Spec vs Schema:

- Protocol = process flow ("Tasks move to_execute/ → completed/")
- Spec = interface contract ("accepts X params, returns Y")
- Schema = file format ("has these fields with these types")

Playbook vs Instruction:

- Playbook = platform-agnostic ("check queues, review stale tasks")
- Instruction = platform-specific ("Claude: use Desktop Commander...")

Directory Structure

```
ai_general/docs/
├── 10_architecture/      WHY
├── 20_registries/        WHAT EXISTS
├── 30_protocols/         HOW IT WORKS (process)
├── 40_specs/             HOW IT WORKS (interface)
├── 50_schemas/           HOW IT WORKS (structure)
├── 60_playbooks/          WHAT TO DO
├── 70_instructions/      HOW TO DO IT
└── 80_quickref/           CHEAT SHEETS
```

5. MCP Servers & Tools Reference

Server Inventory

Desktop Commander

Filesystem operations, process management, file search.

Category	Functions
File I/O	read_file, write_file, write_pdf, edit_block
Directory	list_directory, create_directory, move_file
Search	start_search, get_more_search_results
Processes	start_process, interact_with_process

Codex MCP

Synchronous AI execution (30-60s timeout). NOT a worker.

Tool	Description
codex:codex	Start new Codex session
codex:codex-reply	Continue existing conversation

CLI Agent MCP

Launch and manage CLI agents with role-based bootstrapping.

Tool	Description
launch_agent	Generic launcher (platform + role)
launch_librarian	Memory system curator
launch_dev_lead	Development coordinator
launch_custodian	Repository maintainer
launch_ops	Task execution coordinator
launch_peer_review	Code/design reviewer
launch_tester	Testing and validation
launch_researcher	Corpus search orchestrator
launch_validator	Adversarial cross-checker
kill, attach, send_keys	Session management
list_sessions, get_status	Monitoring

Task Coordination MCP

Playbook-based orchestration and task lifecycle.

Tool	Description
list_playbooks	Available orchestration patterns
get_playbook	Full playbook definition
start_playbook	Create initial task
stop_playbook	Cancel running playbook
list_templates	Available task templates
gen_task	Generate from template
list_tasks	List tasks (filter by platform/status)
get_task	Get task content by ID
move_task	Change task status
list_platforms	Available execution platforms

Knowledge Search MCP

Dual-mode search over 4+ year conversation archive.

Tool	Description
search	4-layer cascade (SEARCH or ANSWER mode)
grep_search	Regex over full content
stats	Knowledge base statistics

Chat Pipeline MCP

Processing pipeline for chat history.

Tool	Description
pipeline_status	Stage counts and status
normalize	Convert raw exports to YAML
chunk_file	Chunk a conversation file
prepare_for_condensation	Full pipeline run
condense_history	Condense a history file
review_quarantine	List failed files

Messages MCP

Inter-AI messaging — broadcasts and direct messages.

Tool	Description
broadcast	Send to all agents
send_direct	Send to specific agent
list_broadcasts	Recent broadcast messages
list_direct	Inbox for a recipient
check_responses	Responses to a message
acknowledge	Mark message read

Prompting MCP

Deliver prompts to AI targets with timing and verification.

Tool	Description
send_prompt	Send to any AI target
is_busy	Check if target is processing
list_sessions	Active tmux sessions
observe_session	Capture CLI session state
wait_response	Wait for response pattern
send_to_session	Low-level tmux send

Todo MCP

Todo/task management with kanban views and status tracking.

Tool	Description
list	List todos with filtering
get	Get details by ID or reference
create	Create new todo
set_status	Update status
add_flag / remove_flag	Manage flags
add_tag / remove_tag	Manage tags
complete	Mark done, move to completed/
trash	Soft delete

kanban	Kanban board view
--------	-------------------

Chrome Control MCP

Browser tab management and JavaScript execution.

Tool	Description
open_url	Open URL in Chrome
list_tabs, get_current_tab	Tab information
switch_to_tab, close_tab	Tab management
execute_javascript	Run JS in tab
get_page_content	Extract page text

Chat Tools MCP

Chat navigation, export, import, and browser interaction.

Tool	Description
open_chat, open_project	Navigation
new_chat	Open new chat
get_messages	Extract messages
send_message	Send to Claude
export_chats	Export to JSON
import_chat	Export → convert → condense
continue_in_new_chat	Fork with condensed context

Tool Usage by Actor

MCP Server	Desktop Claude	CLI Agents	Codex MCP
Desktop Commander	✓	✓	✓
Codex MCP	✓	✓ (potential)	—
CLI Agent	✓	✓	—
Task Coordination	✓	✓	—
Knowledge Search	✓	✓	—
Chat Pipeline	✓	✓	—
Messages	✓	✓	—

Prompting	✓	✓	—
Todo	✓	✓	—
Chat Tools	✓	—	—
Chrome Control	✓	—	—

6. Glossary & Knowledge Index

This glossary solves the bootstrap problem: AIs need to know what's IN files to know WHEN to load them. It provides term recognition without requiring full document loading.

Architecture Terms

Term	Definition	Reference
AI Root	Root directory ~/Documents/AI/ai_root/	architecture_overview
Layer Model	3-tier: communication (ai_comms), implementation (ai_claude), knowledge (ai_memories)	architectural_layer_model
Bootstrap Problem	Need to know file contents to know when to load them — solved by this glossary	schema_knowledge_glossary
Bootstrap Hierarchy	CLI load order: global.md → platform/.md → role/role.yml → tasking.md	ai_general/prompts/

Entities — Platforms

Term	Aliases	Definition
Desktop Claude	Desktop, Claude	Primary orchestrator in desktop app — coordination, memory, decisions. NOT a CLI
Claude CLI	Claude Code, claude_cli	Claude in terminal via claude_cli.py. Tmux sessions, agent profiles (-A), auto mode (-a)
Codex CLI	Codex, codex_cli	OpenAI Codex in terminal. Different from Codex MCP (tool vs worker)
Gemini CLI	Gemini Code, gemini_cli	Google Gemini in terminal. Wave orchestration capable. 1M token context
Cline CLI	Cline, cline_cli	Local agentic AI via Cline + llama-server. Qwen3-Coder on port 8081
llama-server	local LLM	llama.cpp inference server on port 8081. Backend for Cline and mcp-openai
ChatGPT	Chatty	Peer AI collaborator. Project shared memory. Alternative perspective

Roles

Role	Aliases	Definition	Scope
Librarian	memory librarian	Memory system curator, chat history pipeline	ai_memories/

Dev Lead	dev lead, development coordinator	Development coordinator, owns todos and task creation	ai_general/todos/
Custodian	repo custodian	Repository maintainer, structural integrity, hygiene	ai_root/
Ops	operations	CLI coordination and task execution, dispatches agents	ai_comms/
Peer Review	peer reviewer	Code/design reviewer, quality assurance	cross-cutting
Tester	qa tester	Testing specialist, validation and verification	cross-cutting

Tools

Term	Aliases	Definition
Codex MCP	Codex server, codex tool	OpenAI Codex as MCP Server — synchronous tool, NOT a worker. 30-60s timeout
CLI-Agent MCP	cli-agent MCP	MCP server for launching and managing CLI agents
Desktop Skills	claude.ai skills, SKILL.md	Model-invoked instruction packages for Claude Desktop. Auto-activate based on context
Claude Code Plugins	plugins, /plugin	Extension system for Claude Code — commands, agents, skills, MCP servers
Plugin Marketplace	marketplace	Git repos hosting plugin catalogs. Add via <code>claude plugin marketplace add</code>
Superpowers	obra superpowers	Core skills library by Jesse Vincent — TDD, debugging, /brainstorm, git-worktrees
Episodic Memory	episodic-memory	Semantic search across past Claude Code conversations

Search System

Term	Aliases	Definition
Search Agent	knowledge search	Gemini-based dual-mode search over 4yr conversation archive
Search Cascade	4-layer search	L1 topics → L2 topics+synonyms → L3 content → L4 content+synonyms. Never stop at L1
Answer Mode	editorial mode	Search agent mode for questions — synthesizes editorial with [N] citations
Search Mode	results mode	Search agent mode for discovery — curated list with relevance notes

Protocols & Concepts

Term	Aliases	Definition
AT Self-Wake	AT alarm, self-wake	Desktop Claude schedules AT jobs to self-prompt. AT = Desktop's alarm, not CLI's doorbell
Flag Files	state flags, lifecycle flags	Zero-byte files marking task state: *_started, *_completed, *_error, *_cancelled
Claim Prefix	claimed prefix	DEPRECATED v9.0 — replaced by flag files
Message Inserts	memory inserts, INSERT blocks	Structured <<<INSERT>>> blocks for cross-platform persistence markers
Reference Pointers	REF:, file pointers	REF:path/to/file.yml syntax for lazy loading, reduced context
Context Files	context_files list	REF: pointers in role.yml defining files to load for that role
role.yml	role definition	Role config with context_files, duties, ownership. At prompts/roles/{role}/
Time Loop	UI reset	Claude Desktop phenomenon — response fails mid-generation, UI resets context

Workflows

Term	Aliases	Definition
Condensed Chat History	condensed chat	Semantic summary created ON-DEMAND, not pre-stored. ~80% reduction
Chat Recovery	broken chat recovery	Auto-recovery from context exhaustion — detect, export, condense, continue
Task Lifecycle	task states	staged → to_execute → in_progress → completed/error/Cancelled
Task ID Counter	NextId, next_id.sh	Atomic hierarchical ID via mkdir spinlock

Memory System

Term	Definition
Memory Slots	30 slots × ~200 chars each. Hold pointers to external files, not full content
Meta Slots (01-02)	Self-describing: 01=ecosystem_manifest (federation), 02=manifest (structure)

Content Slots (03-10)	Owned memory: user_model, communication, tools, context, limitations, learnings, cross_ai, history
Reserved Slots (11-27)	Available for expansion as memory needs grow
Inbox Slots (28-30)	Web/iOS Claude writes here; Desktop migrates to proper slots
Memory Files	Full-depth storage: user_model.yml, communication_patterns.yml, tool_discoveries.yml, etc.
Manifest	Index file (manifest.yml) defining what each slot is for — slot 02 points here
Ecosystem Manifest	Cross-AI federation registry — slot 01 points here
Federated Memory	Per-AI owned memory with cross-reference capability via ai_ecosystem_manifest.yml
Hot/Warm/Cold	Memory tiers: Hot=loaded, Warm=on-demand, Cold=search/retrieval
Write-Immediate	Discipline: log observations as they occur, never batch or checkpoint

Quick Lookup by Domain

Domain	Key Terms
Architecture	AI Root, Layer Model, Bootstrap Problem
Platforms	Desktop Claude, Claude CLI, Codex CLI, Gemini CLI, Cline CLI
Roles	Librarian, Dev Lead, Custodian, Ops, Peer Review, Tester
Tools	Codex MCP, CLI-Agent MCP, Desktop Skills, Plugins
Search	Search Agent, Search Cascade, Answer Mode, Search Mode
Protocols	AT Self-Wake, Flag Files, Message Inserts, REF Pointers
Memory	Meta Slots, Content Slots, Inbox Slots, Memory Files, Federated Memory
Workflows	Condensed Chat, Chat Recovery, Task Lifecycle

Generated February 7, 2026 — Version 1.2.0